

# A small discussion on PREFIX SUM

Luca Tornatore - I.N.A.F.




DATA SCIENCE &  
ARTIFICIAL INTELLIGENCE



SCIENTIFIC &  
DATA-INTENSIVE COMPUTING

**2024-2025 @ Università di Trieste**




# Algorithm analysis

In the exercise we have introduced the prefix sum; a brief search may convince you that is a fundamental algorithm upon which many other algorithms are built.

The serial algorithm is quite obvious, although the most obvious implementation may be optimized for modern architectures.

In these few slides, I intend to discuss two different parallel implementations. I aim to convince you that not all parallelisms are born equal.

In the following,  **$N$**  is the size of the array and  **$n$**  indicates the number of threads.



# Algorithm I




Let's say that this is our array, of which we want to compute the prefix-sum.  
An obvious *phase-1* is to divide the array in chunks and compute the partial prefix sum for each chunk



The first chunk is correct, while the others lack the contributions from the previous chunks







# Algorithm I

We can solve this by a recurrent subdivision of the work among *all* the tasks, with  $n-1$  steps


Step 1




Step 2



Step  $n-1$





# Algorithm I

How many ops are performed in this algorithm ?

Step 0 :  $N/n \times n - n$  operations ( the term  $-n$  is due to the fact that the first elements of each chunk are untouched in this step



Step 1 :  $N - N/n$  operations




Step 2 :  $N - 2 \times N/n$  operations



Step  $n-1$  :  $N - (n-1) \times N/n$  operations






# Algorithm I

How many ops are performed in this algorithm ?

We can express that recurrence as

$$\begin{aligned} N + \left(N - \frac{N}{n}\right) + \left(N - 2 \times \frac{N}{n}\right) + \dots + \left(N - (n-1) \times \frac{N}{n}\right) - n = \\ \sum_{i=0}^{i=n-1} \left(N - i \times \frac{N}{n}\right) - n = \\ nN - \frac{N}{n} \times \frac{n(n-1)}{2} - n = \\ nN - \frac{Nn}{2} + \frac{N}{2} - n = \\ \frac{N}{2}(n+1) \end{aligned}$$



# Algorithm II




Let's start from the first *phirst* phase as before



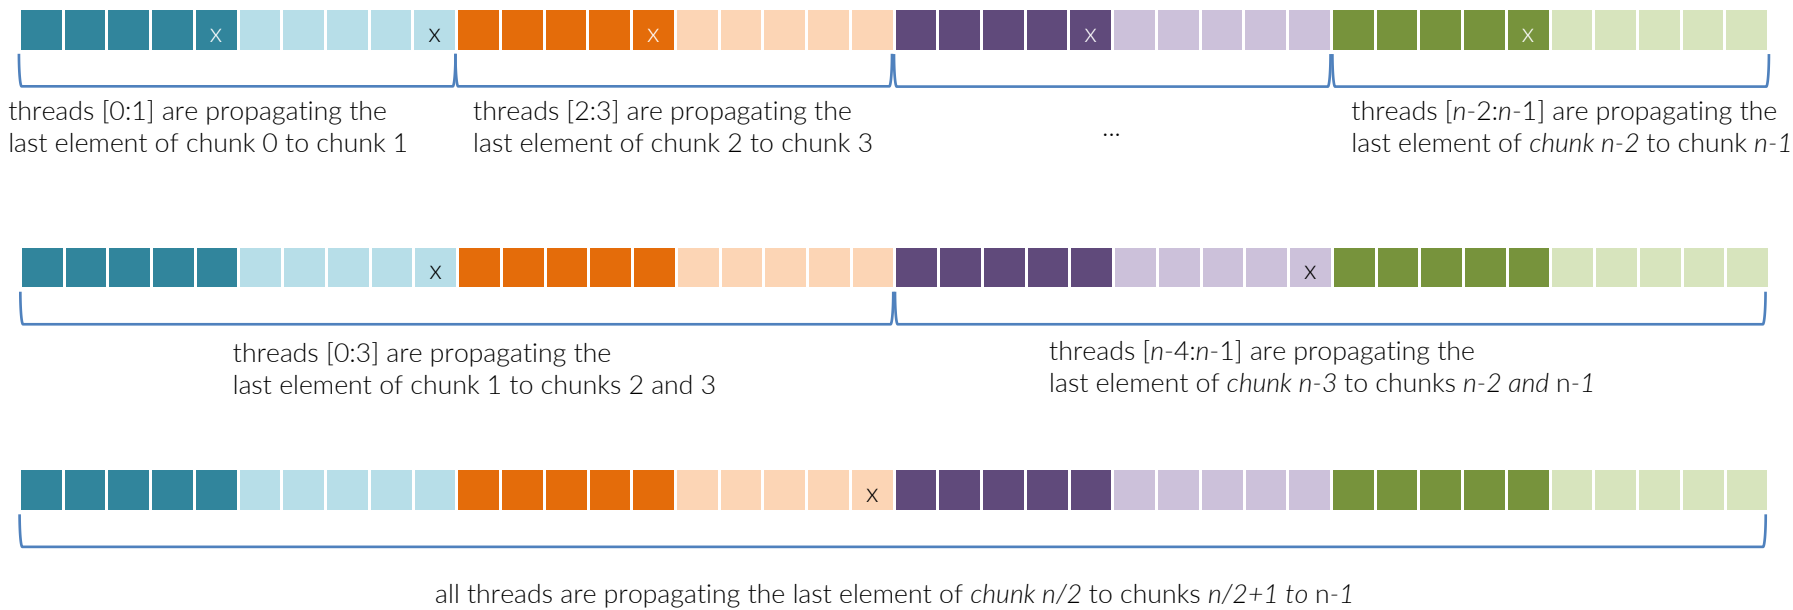
The first chunk is correct, while the others lack the contributions from the previous chunks






# Algorithm II

Let's design a different recurrent subdivision of the work among *all* the workers (we're using 8 threads instead of 4 to illustrate it)



*Note: you can easily adapt it to work with a non-power-of-two number of threads*





# Algorithm II

How many ops are performed in this algorithm ?

In the step 0 it performs  $N-n$  operations.

In step 1 it performs  $N/n \times n/2$


In step  $i$  it performs  $j \times \frac{N}{n} \times \frac{n}{2j}$  where  $j = i^2$

Since there will be  $k = \log_2 n$  steps, we can sum up as

$$N - n + \sum_{i=1}^k \left( \frac{N}{2} \right) =$$

$$N - n + \frac{N}{2} k =$$


$$N \left( 1 + \frac{k}{2} \right) - n = \mathbf{N \left( 1 + \frac{\log_2 n}{2} \right) - n}$$




# Algorithm I vs II

Algorithm II is manifestly superior:  
the ratio of the ops number goes  
as (ignoring the  $-n$  terms)

$$R = \frac{N/2(n+1)}{N \left(1 + \frac{\log_2 n}{2}\right)} \approx \frac{n}{\log_2 n}$$





# | Algorithm I vs II

Even more important is the fact that the speed-up of the two algorithms (supposing that the run time goes a  $\#ops/n$ ) has very different behaviour

algorithm 1:

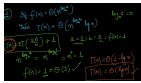
$$\text{run-time} = \frac{N}{2} \times \frac{n+1}{n}$$

$$\text{speed-up} = T_s(N) / T_p(N,n) = \frac{2n}{n+1} \approx 2$$

algorithm 2:

$$\text{run-time} = \frac{N}{n} \times \left(1 + \frac{\log_2 n}{2}\right)$$

$$\text{speed-up} = \frac{n}{1+\log_2 n} \approx \frac{n}{\log_2 n}$$


$$A_j = A_j + A_{j-1} \quad \text{for } j = 1 \text{ to } n$$

Prefix  
Sum

# Memory issues

There may be another (severe) limiting factor to the *scalability* (i.e. the speed-up growth when you increase the number of threads) that is linked to the memory access.

We'll discuss this matter specifically when we'll discuss the threads-memory affinity.